

# The Modern PyTorch Workflow: From Boilerplate to Lightning.

Scaling PyTorch Experiments Everywhere with Lightning 2.x + Fabric.

Siman Giri

# Disclaimer!!!

- I am not a core **PyTorch developer** – I haven't been involved in designing its internal packages.
- Then Why am I here:
  - While I didn't design PyTorch, I've been living in its world for over four years doing AI research.
  - Today, I'll share my lessons learned,
    - plus, the latest Lightning 2.x and Fabric updates that
      - help you scale models without losing your sanity.

# 1. Introduction.

# About Me!!!

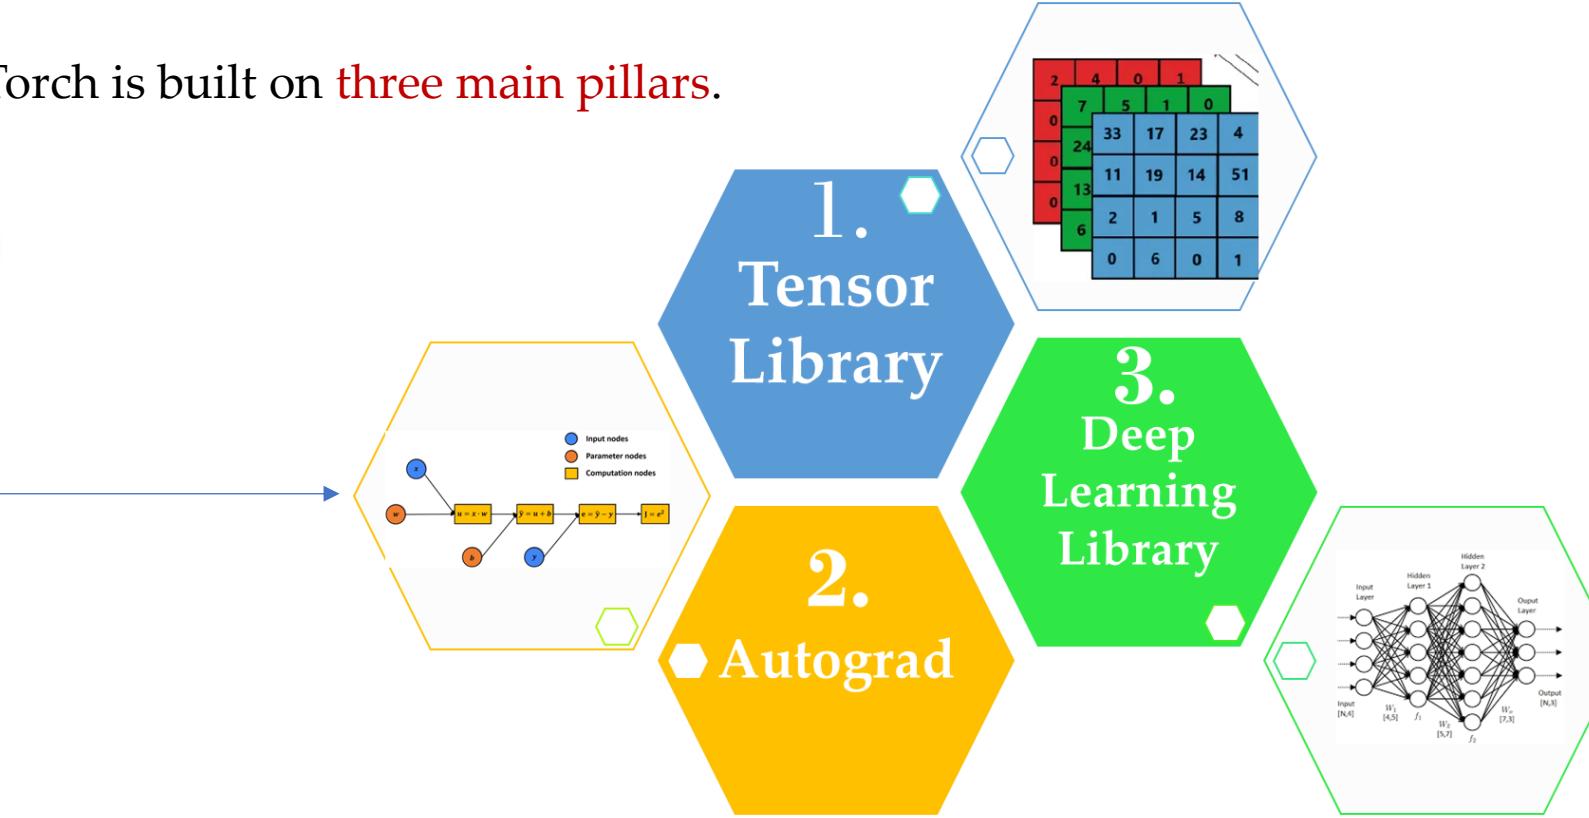
- **Siman Giri**
  - **B.E. Biomedical, MS Applied Mathematics and Machine Learning, MSc Mathematics**
    - **Doctoral Researcher @ Madan Bhandari University of Science and Technology.**
      - **Advised By: Dr. Suresh Manandhar**
      - Research Focus: Theoretical and Mathematical Framework for Reasoning in Multimodal AI.
      - Current Work: Robust Object Tracking for small Video Language Model.
    - **Head @ Centre for AI and Emerging Technologies at Herald College Kathmandu.**
      - Research Unit: IoT & Robotics Unit, Vision & Language Research Unit
      - Objective: Prepare Undergrad for world of Modern-Day AI Research
    - **Researcher @ TU Supercomputing Center.**
      - Research Focus: HPC Infrastructure for AI and Quantum Research.
      - **Advised By: Dr. Madhav Pd. Ghimire (Head of the Center).**

# 1.1 Understanding PyTorch: The Three Pillars.

- PyTorch is built on **three main pillars**.



PyTorch

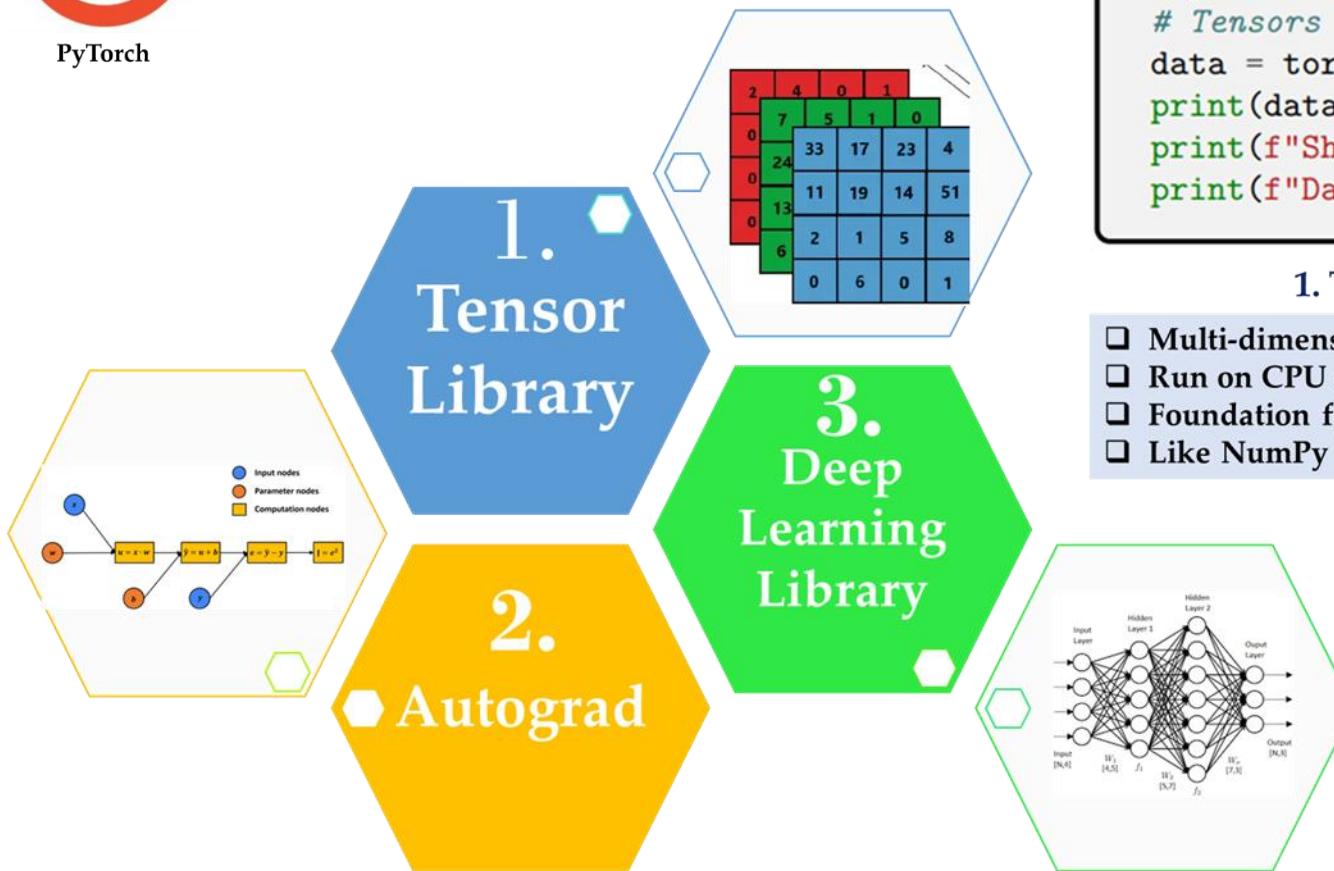


- ✓ At its core, PyTorch combines a **flexible tensor library**, a dynamic **automatic differentiation system**, and **practical deep learning utilities**.
- ✓ Together, these components enable researchers to **efficiently build, optimize, and train deep neural network models**.

# 1.1 Understanding PyTorch: The Three Pillars.



PyTorch



## PyTorch Tensor Example

```
import torch
# Tensors are like NumPy arrays but with superpowers
data = torch.tensor([[1, 2], [3, 4]])
print(data)
print(f"Shape: {data.shape}")
print(f"Data type: {data.dtype}")
```

### 1. Tensors: The Universal Data Container

- Multi-dimensional arrays that can hold any type of data (images, text, audio).
- Run on CPU or GPU seamlessly → massive speedup for computations.
- Foundation for everything: All data in PyTorch flows as tensors through the network.
- Like NumPy but better: Similar interface but with GPU support and gradient tracking.

# 1.1 Understanding PyTorch: The Three Pillars.



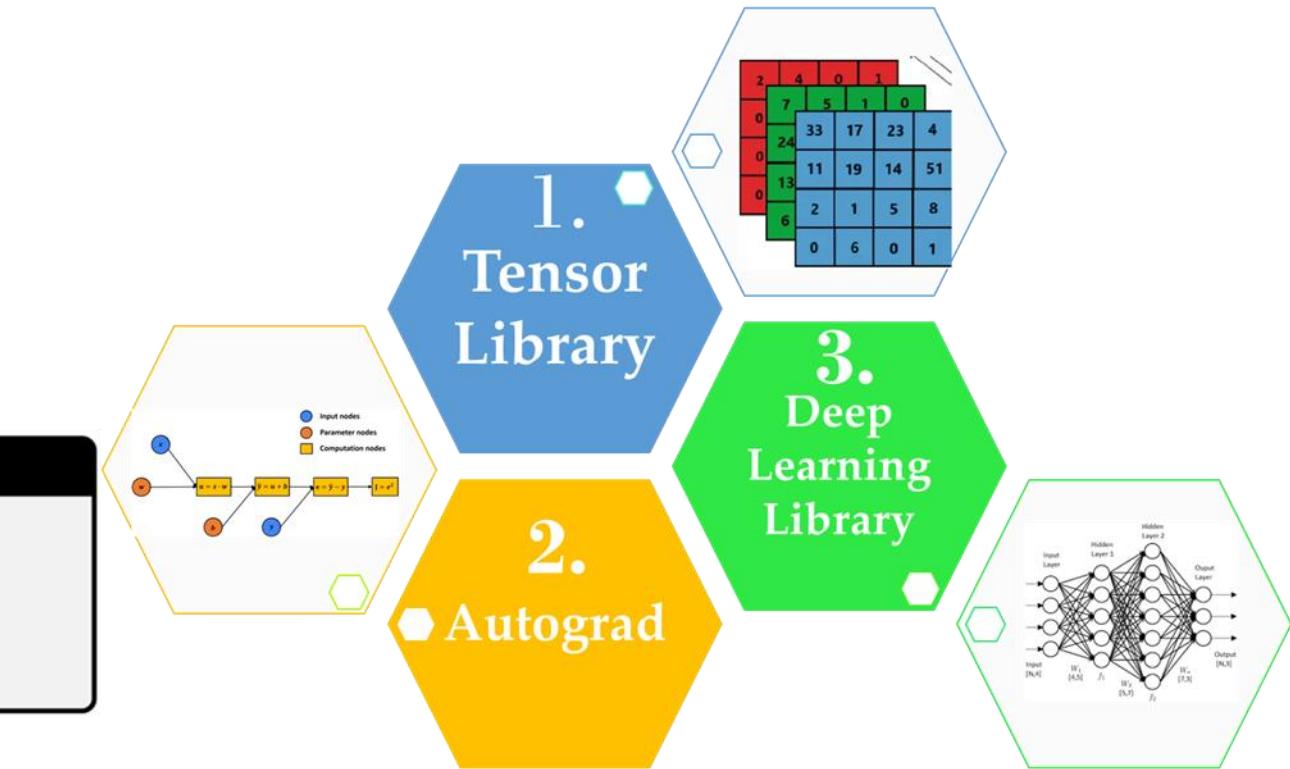
PyTorch

## Autograd Example

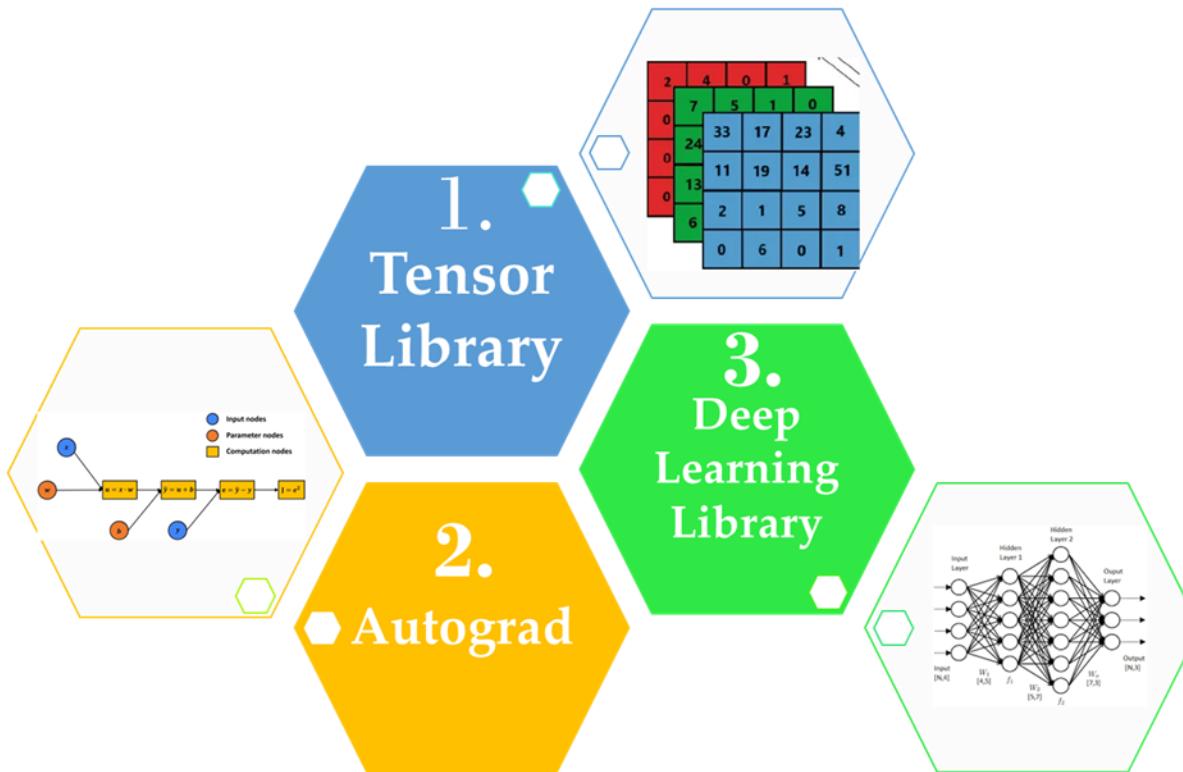
```
x = torch.tensor([2.0], requires_grad=True)
y = x ** 2 # y = x2
y.backward() # Auto-computes dy/dx = 2x
print(x.grad) # Output: tensor([4.]) + Magic!
```

## 2. Autograd : Automatic Differentiation Engine

- ❑ Automatically computes gradients (derivatives) for optimization
- ❑ Builds computation graphs behind the scenes
- ❑ Enables backpropagation with just .backward() call
- ❑ Why it matters: Manually calculating derivatives for deep networks (millions of parameters) is impossible;
  - ❑ autograd does it automatically.



# 1.1 Understanding PyTorch: The Three Pillars.



## Deep Learning Libraries

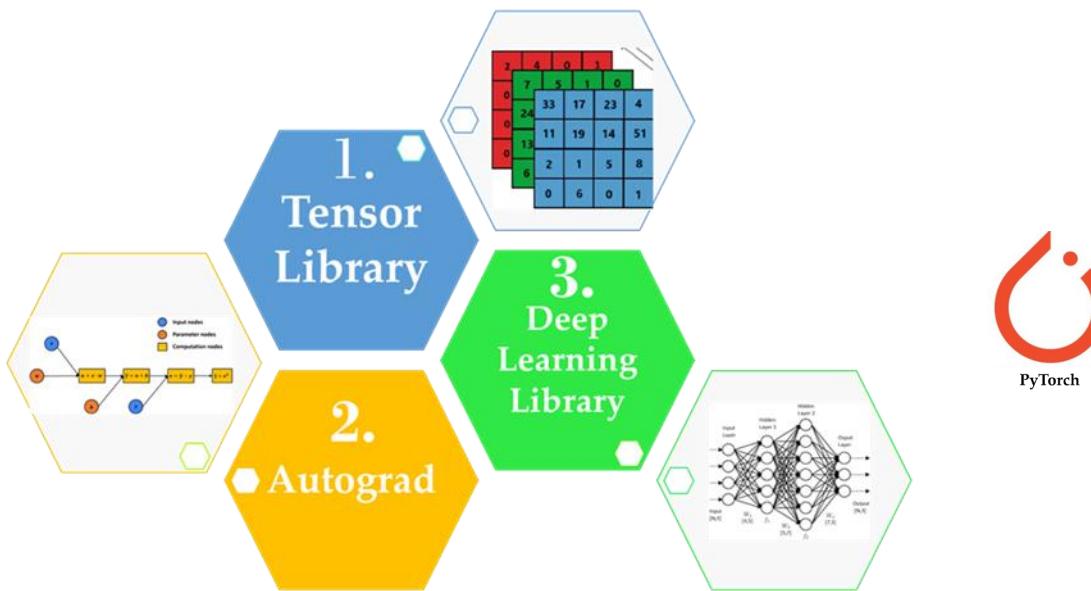
```
import torch.nn as nn

# Pre-built layers
model = nn.Sequential(
    nn.Linear(100, 50), # Fully connected layer
    nn.ReLU(),           # Activation function
    nn.Linear(50, 10)   # Output layer
)
```

## 3. Neural Network Library: Pre-built AI Components

- **Pre-implemented layers** (convolutional, recurrent, attention, etc.)
- **Loss functions** (cross-entropy, MSE, etc.)
- **Optimizers** (SGD, Adam, etc.)
- **Utilities** (data loaders, transformers).

# 1.2 Why These Matters?



- ✓ Tensors → represent data.
- ✓ Autograd → computes gradients.
- ✓ Deep learning libs → build, train, and optimize models.

## How They Work Together: A Complete Picture

```
# 1. TENSORS: Create input data
inputs = torch.randn(32, 3, 224, 224) # Batch of 32 images

# 2. NEURAL NETWORK: Define model
model = torchvision.models.resnet18()

# 3. AUTGRAD: Training loop
outputs = model(inputs)
loss = nn.CrossEntropyLoss()(outputs, labels)
loss.backward() # Autograd computes all gradients!
optimizer.step() # Update weights using those gradients
```

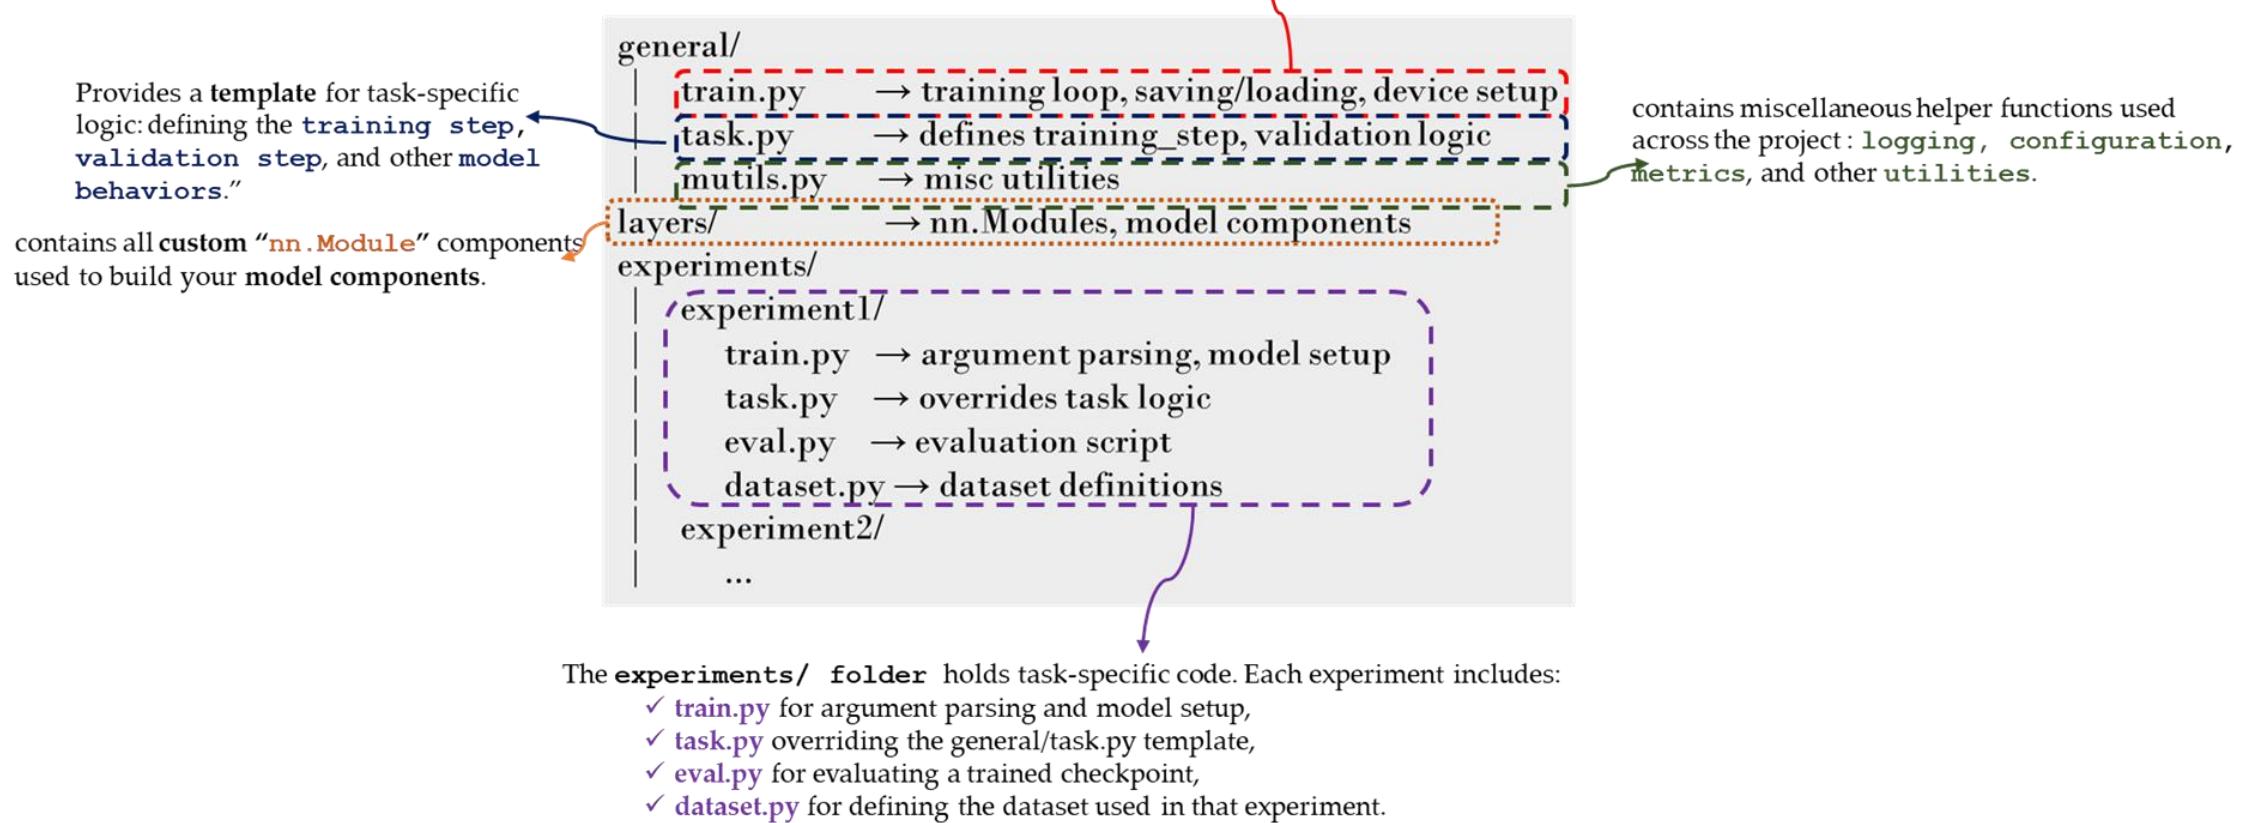
- ✓ Together they form a **flexible, low-level deep learning framework** that researchers love.

# 1.2 The PyTorch Paradox: Unlimited Freedom, Repeated Structure.

- One possible way to Build Projects using PyTorch – {**Suggestive**} Only.

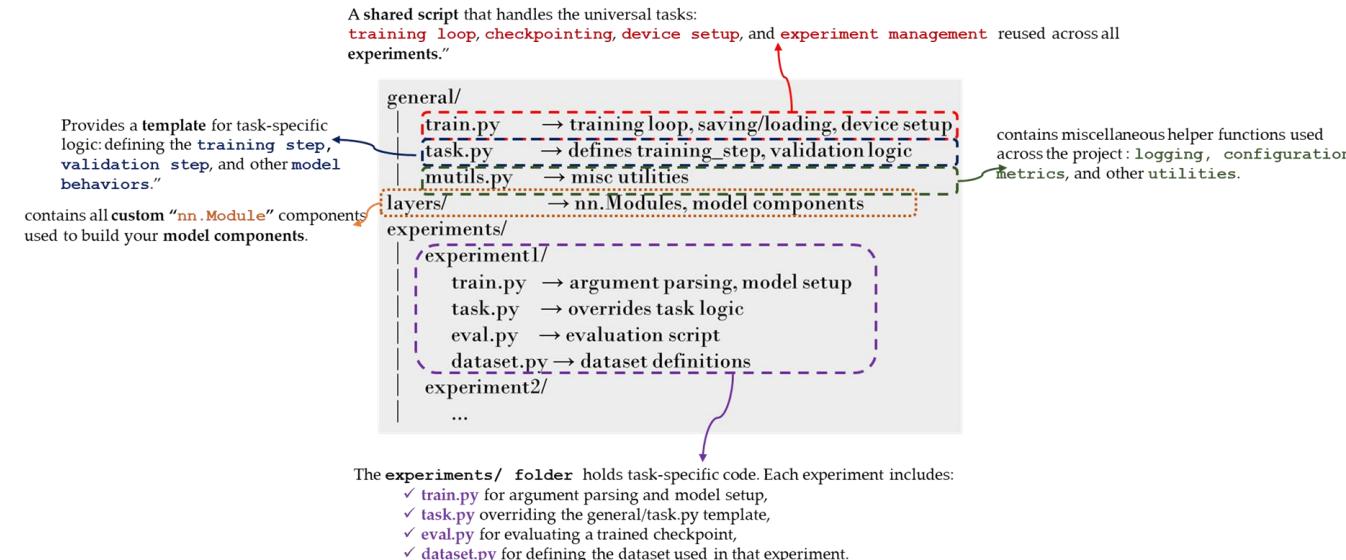
A shared script that handles the universal tasks:

`training loop, checkpointing, device setup, and experiment management` reused across all experiments.”



# 1.2 The PyTorch Paradox: Unlimited Freedom, Repeated Structure.

- One possible way to Build Projects using PyTorch – {**Suggestive**} Only.



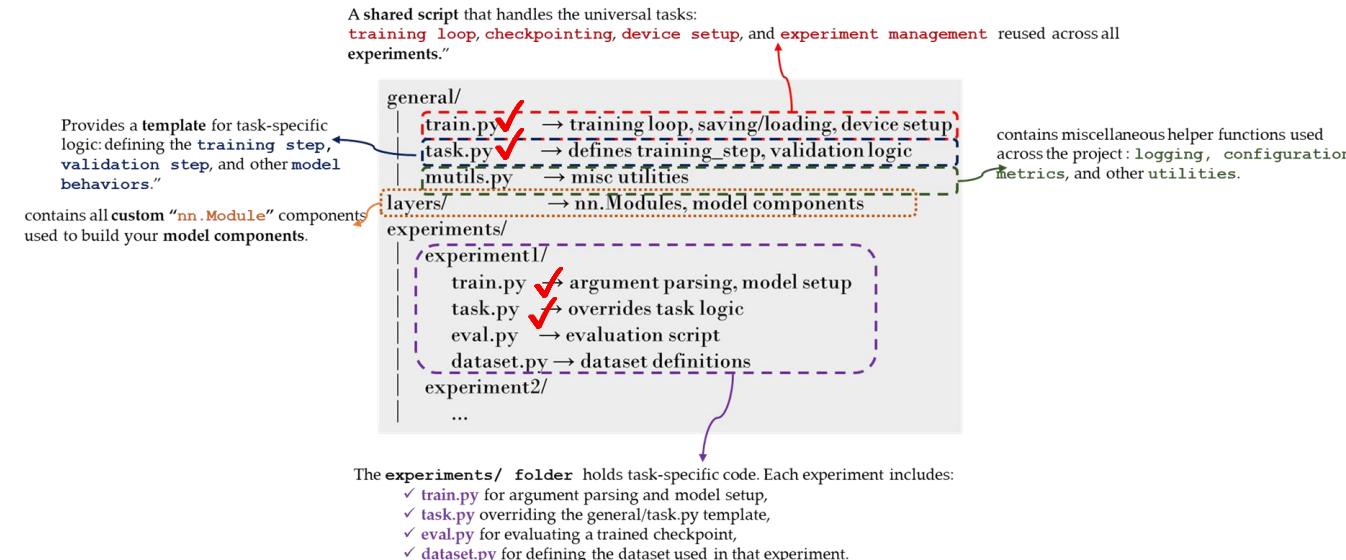
## □ PyTorch's Philosophy:

- ✓ **Explicit over implicit:**
  - ✓ You see and control every step
- ✓ **Pythonic:**
  - ✓ It feels like regular Python code
- ✓ **Dynamic:**
  - ✓ Computation graphs are built on-the-fly

If PyTorch is great, then **what** are the limitations?

# 1.2 The PyTorch Paradox: Unlimited Freedom, Repeated Structure.

- One possible way to Build Projects using PyTorch – {**Suggestive**} Only.



## □ PyTorch's Philosophy:

- ✓ **Explicit over implicit:**
  - ✓ You see and control every step
- ✓ **Pythonic:**
  - ✓ It feels like regular Python code
- ✓ **Dynamic:**
  - ✓ Computation graphs are built on-the-fly

If PyTorch is great, then **what** are the limitations?

# 1.3 PyTorch: Flexibility with a Side of Boilerplate

The Universal PyTorch Training Template

```
# THE BOILERPLATE YOU WRITE 100 TIMES
def train_one_epoch(model, dataloader, optimizer, device):
    model.train()
    for batch in dataloader:
        x, y = batch
        x, y = x.to(device), y.to(device)    # ← Always this line

        optimizer.zero_grad()
        y_hat = model(x)
        loss = loss_fn(y_hat, y)
        loss.backward()
        optimizer.step()

        # Logging (different each time)
        writer.add_scalar('loss', loss, step)

        # Checkpoint logic (reinvented each project)
        if step % 1000 == 0:
            torch.save(model.state_dict(), f'ckpt_{step}.pt')

    return average_loss
```

Fig: This is what a Standard Training Loop in PyTorch Looks like, So what is the challenge?

# 1.3 PyTorch: Flexibility with a Side of Boilerplate

The Universal PyTorch Training Template

```
# THE BOILERPLATE YOU WRITE 100 TIMES
def train_one_epoch(model, dataloader, optimizer, device):
    model.train()
    for batch in dataloader:
        x, y = batch
        x, y = x.to(device), y.to(device) # ← Always this line

        optimizer.zero_grad()
        y_hat = model(x)
        loss = loss_fn(y_hat, y)
        loss.backward()
        optimizer.step()

        # Logging (different each time)
        writer.add_scalar('loss', loss, step)

        # Checkpoint logic (reinvented each project)
        if step % 1000 == 0:
            torch.save(model.state_dict(), f'ckpt_{step}.pt')

    return average_loss
```



## Experiment A

- ❑ `train_a.py`
- ✓ custom args
- ✓ unique logs
- ✓ special ckpt
- ✓ Manual DDP

## Experiment B

- ❑ `train_b.py`
- ✓ custom args
- ✓ unique logs
- ✓ special ckpt
- ✓ Manual DDP

## Experiment C

- ❑ `train_a.py`
- ✓ custom args
- ✓ unique logs
- ✓ special ckpt
- ✓ Manual DDP



THE SAME BOILERPLATE, REPEATED & MODIFIED  
(training loop, device mgmt., logging, etc.)

**The irony:** We write the same **infrastructure code** repeatedly while trying to do novel research.

# 1.4 The Hidden Costs of Reinventing the Wheel.

## 1. Repetitive Boilerplate

- **Problem:**

- Duplicated code across experiments increases maintenance cost.
- Small mistakes in copy – pasted loops lead to silent training bugs.
- Researchers spend more time re – writing loops than innovating on models.

- **Impact:**

- Slows down experimentation, reduces productivity, and creates inconsistencies across projects.

## 2. Manual Device Handling

- **Problem:**

- Device placement (`.to(device)`) scattered throughout code leads to errors.
- Switching from **CPU** → **GPU** → **Multi-GPU** → **TPU** forces complex code rewrites.
- Device logic **pollutes** model/training logic.

- **Impact:**

- Fragile workflows, increased debugging time, and **non – portable code** across environments.

## 3. DIY Experiment Management

- **Problem:**

- Researchers must implement **logging**, **checkpointing**, **reproducibility**, and **monitoring** on their own.
- Each **implementation** is different, inconsistent, and prone to errors.
- Reproducibility relies on **manually** setting seeds and managing randomness – often overlooked.

- **Impact:**

- Harder collaboration, poorer reproducibility, and duplicated engineering effort across teams.

## 4. Scaling Complexity

- **Problem:**

- Distributed training (**DDP**, **FSDP**, **TPUs**, **AMP**) requires deep engineering knowledge.
- Wrong configuration **silently produces** incorrect gradients or broken synchronization.
- Scaling even a simple model **requires rewriting** loops and setup code.

- **Impact:**

- High barrier to scaling, **greater likelihood** of distributed **bugs**, and slower research iteration.

# 1.4.1 The Hidden Costs of Reinventing the Wheel.

## 5. Dataset Duplication

- **Problem:**

- Slight changes in augmentation or data splits lead to almost identical `dataset.py` files.
- Each experiment contains repeated `DataLoader` logic.
- Hard to maintain consistency across experiments.

- **Impact:**

- Code bloat, inconsistent preprocessing, and increased chance of data-related bugs.

## 6. Error-Prone Infrastructure

- **Problem:**

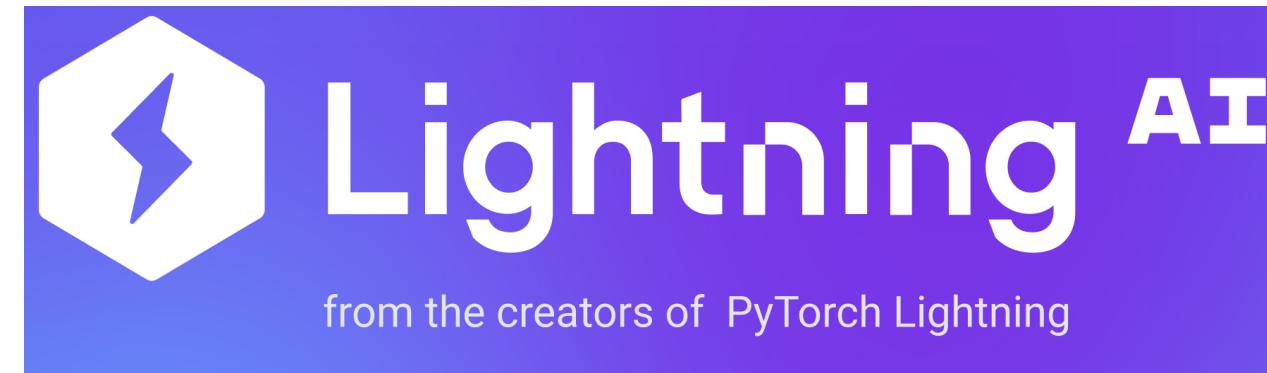
- Custom training loops hide subtle bugs (gradient accumulation errors, `missing .eval()`, incorrect `zero_grad`, wrong validation frequency).
- Distributed training bugs are extremely hard to detect.
- Checkpointing and resuming training is fragile and usually incomplete.

- **Impact:**

- Unreliable experiments, wasted compute, and difficulty reproducing results.

❑ Raw PyTorch gives full control → but with control comes the burden of writing and maintaining training infrastructure.

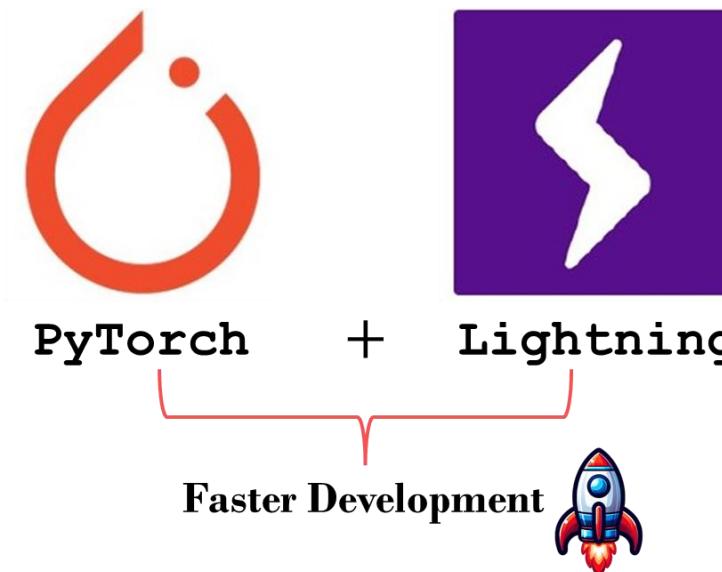
❑ "What if we could keep PyTorch's flexibility but eliminate this repetition?"



## 2. Scale, Don't Struggle: Modern Distributed PyTorch with Lightning & Fabric

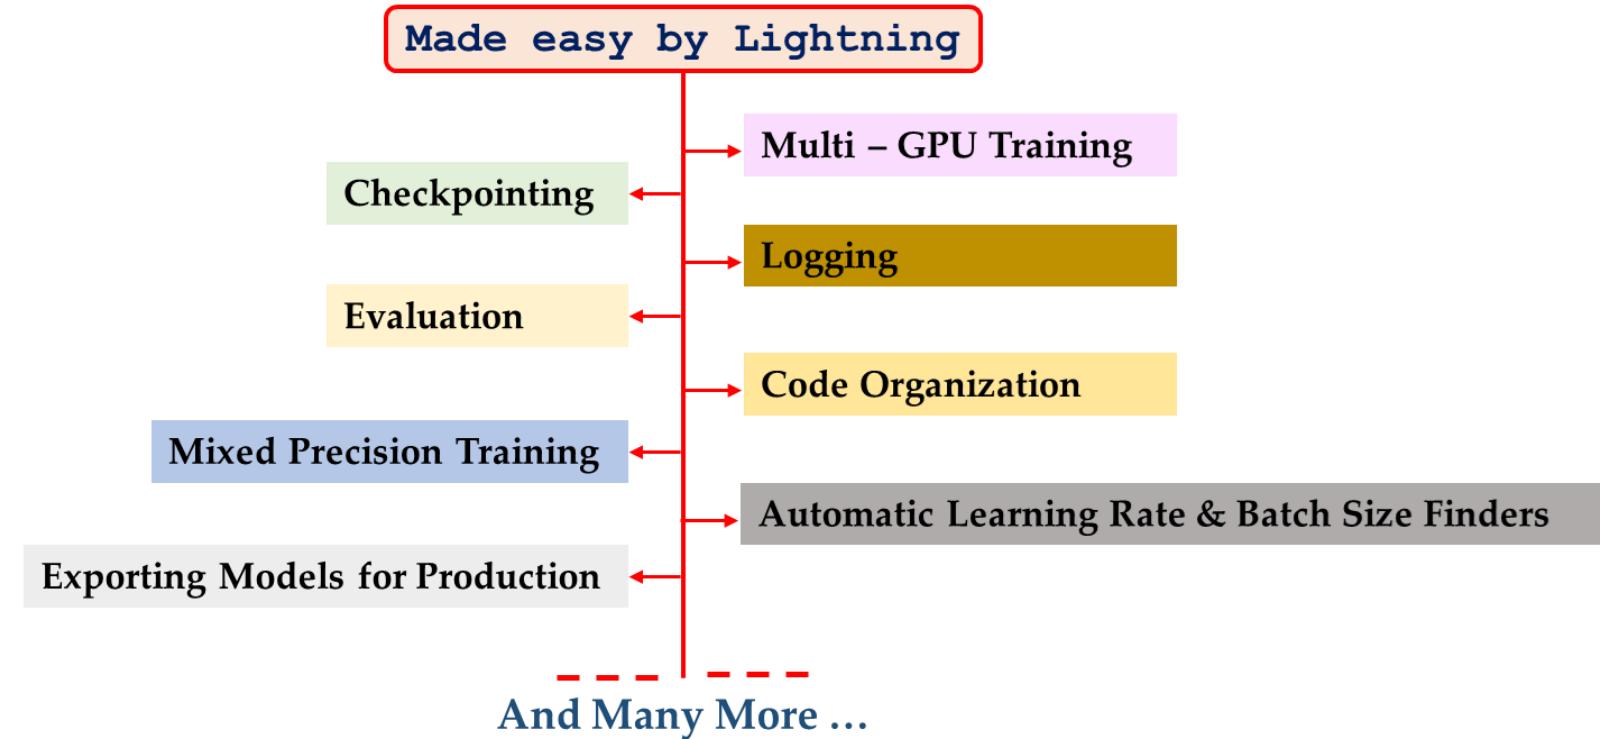
## 2.1 Introducing {PyTorch} Lightning Because I'd Rather Debug My Ideas, Not My Training Loop.

- PyTorch Lightning is a **high-level deep learning framework** built on top of PyTorch.
- Its primary goal is simple:
  - *"Let researchers focus on the science, not the boilerplate."*
- Lightning **separates research logic from engineering code**, giving you a structured, clean workflow for experiments
  - without forcing you to abandon **PyTorch's flexibility**.





## 2.2 Introducing {PyTorch} Lightning: Less Code, More Results: Lightning in Action.



**Big Claim:** “By adopting PyTorch Lightning you could reduce Engineering Overhead in PyTorch Workflows.”



## 2.3 Lightning by Example: Clean Code, Same Science.

### 1. Full Flexibility – Use Raw PyTorch where it Matters:

- You still write pure PyTorch inside the model.

#### Full Flexibility

```
import torch
import torch.nn as nn
import pytorch_lightning as pl
class LitClassifier(pl.LightningModule):
    def __init__(self, lr=1e-3):
        super().__init__()
        self.save_hyperparameters()
        self.model = nn.Sequential(
            nn.Linear(784, 256),
            nn.ReLU(),
            nn.Linear(256, 10)
        )
        self.loss_fn = nn.CrossEntropyLoss()

    def forward(self, x):
        return self.model(x)
```



## 2.3 Lightning by Example: Clean Code, Same Science.

### 2. Reproducible + Readable :

- Decoupled Research & Engineering

Reproducible + Readable

```
def training_step(self, batch, batch_idx):  
    x, y = batch  
    logits = self(x)  
    loss = self.loss_fn(logits, y)  
    self.log("train_loss", loss)  
    return loss  
  
def validation_step(self, batch, batch_idx):  
    x, y = batch  
    logits = self(x)  
    acc = (logits.argmax(dim=1) == y).float().mean()  
    self.log("val_acc", acc)
```

Reproducible + Readable

```
def configure_optimizers(self):  
    return torch.optim.Adam(self.parameters(), lr=self.hparams.lr)
```

#### □ Why this proves the point:

- ✓ Research logic = `training_step`, `validation_step`
- ✓ Engineering logic (loops, devices, AMP, DDP) = `Trainer`
- ✓ `self.save_hyperparameters()` → full experiment reproducibility



## 2.3 Lightning by Example: Clean Code, Same Science.

### 3. DataModule – Dataset Handling Without Duplication:

DataModule

```
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
class MNISTDataModule(pl.LightningDataModule):
    def __init__(self, batch_size=64):
        super().__init__()
        self.batch_size = batch_size

    def setup(self, stage=None):
        transform = transforms.ToTensor()
        self.train_ds = datasets.MNIST(".", train=True, download=True,
                                       transform=transform)
        self.val_ds = datasets.MNIST(".", train=False, download=True,
                                       transform=transform)

    def train_dataloader(self):
        return DataLoader(self.train_ds, batch_size=self.batch_size)

    def val_dataloader(self):
        return DataLoader(self.val_ds, batch_size=self.batch_size)
```

Change dataset, augmentation, or batch size → no change to training loop.



## 2.3 Lightning by Example: Clean Code, Same Science.

### 4. Simple Multi-GPU / TPU / Mixed Precision – No Code Changes:

#### Simple Multi-GPU / TPU / Mixed Precision

```
trainer = pl.Trainer(  
    accelerator="gpu",  
    devices=4,      # 1 GPU, 4 GPUs, TPU, HPU / same code  
    precision=16,   # mixed precision  
    max_epochs=10,  
    log_every_n_steps=50  
)  
  
model = LitClassifier()  
data = MNISTDataModule()  
  
trainer.fit(model, datamodule=data)
```

- This replaces hundreds of lines of PyTorch boilerplate:
  - ✓ No `.to(device)`, No DDP setup, No AMP context managers, No manual logging or checkpoint code



## 2.3 Lightning by Example: Clean Code, Same Science.

### 5. Built – in Testing – Battle –Tested Infrastructure:

#### Built-in Testing

```
trainer.test(model, datamodule=data)  
trainer.predict(model, datamodule=data)
```

#### What Lightning already tests for you:

- ✓ Distributed correctness
- ✓ Mixed precision stability
- ✓ Checkpoint restore
- ✓ Gradient accumulation
- ✓ Fault-tolerant training

*"Lightning doesn't change how you think about models – it changes **how much infrastructure** you have to write."*



# PyTorch Lightning

## 3. Organizing Your PyTorch Code with Lightning. {How Lightning organizes your code?}



# 3.1 The Lightning Mental Model

- *Lightning* organizes **PyTorch** code into 3 core components:

Component	Responsibility
LightningModule	What the model learns
DataModule	Where data comes from
Trainer	How training runs

“*Lightning* enforces separation of concerns → models, data, and training infrastructure never mix.”



## 3.2 The LightningModule

- A **LightningModule** organizes your **PyTorch** code into 6 sections:

1.	INITIALIZATION → <code>__init__()</code> & <code>setup()</code>
2.	TRAIN LOOP → <code>training_step()</code>
3.	VALIDATION LOOP → <code>validation_step()</code>
4.	TEST LOOP → <code>test_step()</code>
5.	PREDICTION LOOP → <code>predict_step()</code>
6.	OPTIMIZERS → <code>configure_optimizers()</code>

## 3.2.1 The LightningModule – Initialization.

### Initialization

```
class LitClassifier(pl.LightningModule):
    def __init__(self, num_classes=10, hidden_dim=128):
        super().__init__()
        # Save hyperparameters (auto-logged!)
        self.save_hyperparameters()

        # Define your model architecture (pure PyTorch)
        self.layers = nn.Sequential(
            nn.Linear(28*28, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, num_classes)
        )

    def setup(self, stage=None):
        """Called at beginning of fit/test/predict"""
        # OPTIONAL: Prepare data, initialize weights, etc.
        if stage == "fit":
            # Load pretrained weights, etc.
            pass
```

- What it does:
  - ✓ `__init__`: Define your model architecture, layers, hyperparameters
  - ✓ `setup()`: Optional hook for data-dependent initialization
- Key benefit: `save_hyperparameters()` automatically logs everything

## 3.2.2 The LightningModule – Train Loop.

Train Loop (`training_step()`)

```
def training_step(self, batch, batch_idx):
    # batch comes from your train DataLoader
    x, y = batch

    # 1. Forward pass (no .to(device) needed!)
    y_hat = self(x)  # Calls self.forward()

    # 2. Compute loss
    loss = F.cross_entropy(y_hat, y)

    # 3. Log metrics (auto-logged to all loggers)
    self.log("train_loss", loss,
             on_step=True,  # Log per batch
             on_epoch=True,  # Also log epoch average
             prog_bar=True) # Show in progress bar

    # 4. Return loss (Lightning handles backward() automatically)
    return loss
```

□ What it does:

- ✓ Defines what happens per training batch
- ✓ Lightning automatically calls:
  - `loss.backward()`
  - `optimizer.step()`
  - `optimizer.zero_grad()`

□ Key benefit: Automatic logging with `self.log()`

### 3.2.3 The LightningModule – Validation Loop.

```
Validation Loop (validation_step())

def validation_step(self, batch, batch_idx):
    x, y = batch
    y_hat = self(x)
    loss = F.cross_entropy(y_hat, y)

    # Compute accuracy
    preds = torch.argmax(y_hat, dim=1)
    acc = (preds == y).float().mean()

    # Log validation metrics
    self.log("val_loss", loss, on_epoch=True)
    self.log("val_acc", acc, on_epoch=True)

    # Optional: Return values for custom epoch reduction
    return {"val_loss": loss, "val_acc": acc}

    # OPTIONAL: Custom epoch-level aggregation
    def on_validation_epoch_end(self):
        # Access all validation_step outputs if needed
        avg_loss = torch.stack([x["val_loss"]] ...
            for x in self.validation_step_outputs]).mean()
```

#### □ What it does:

- ✓ Runs **during training** (after each epoch) and **during validation**
- ✓ No gradient computation (automatic `torch.no_grad()`)
- ✓ Metrics logged with `on_epoch=True` for proper averaging

## 3.2.4 The LightningModule – Test Loop.

Test Loop (`test_step()`)

```
def test_step(self, batch, batch_idx):
    x, y = batch
    y_hat = self(x)

    # Compute test metrics
    preds = torch.argmax(y_hat, dim=1)
    acc = (preds == y).float().mean()

    # Log test metrics
    self.log("test_acc", acc)

    # Optional: Save predictions for analysis
    return {"preds": preds, "targets": y}

def on_test_epoch_end(self):
    # Analyze all test results
    all_preds = torch.cat([x["preds"] for x in self.test_step_outputs])
    all_targets = torch.cat([x["targets"] for x in self.test_step_outputs])
    # Compute confusion matrix, F1 score, etc.
```

### □ What it does:

- ✓ Runs **once** after training is complete
- ✓ For final evaluation on unseen test data
- ✓ Identical structure to `validation_step()` but semantically different

## 3.2.5 The LightningModule – Prediction Loop.

### Prediction Loop (predict\_step())

```
def predict_step(self, batch, batch_idx, dataloader_idx=0):
    # Batch could be raw data, no labels
    x = batch # Just features, no labels

    # Forward pass
    y_hat = self(x)

    # Post-processing (softmax, thresholding, etc.)
    probs = F.softmax(y_hat, dim=1)
    preds = torch.argmax(probs, dim=1)

    # Return predictions in any format
    return {"predictions": preds, "probabilities": probs}
```

#### □ What it does:

- ✓ Inference-only mode (no gradients, no logging)
- ✓ For **production inference** or **batch predictions**
- ✓ Handles data **without labels**
- ✓ Returns raw **predictions** for external use



## 3.2.6 The LightningModule – Optimizers.

Optimizers (configure\_optimizers())

```
def configure_optimizers(self):
    # Basic: Single optimizer
    optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)

    # With learning rate scheduler
    optimizer = torch.optim.SGD(self.parameters(), lr=1e-2, momentum=0.9)
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=30)

    # Return format depends on what you need
    return {
        "optimizer": optimizer,
        "lr_scheduler": {
            "scheduler": scheduler,
            "interval": "epoch", # or "step"
            "frequency": 1,
            "monitor": "val_loss", # For ReduceLROnPlateau
        }
    }

    # Multiple optimizers (GANs, meta-learning)
    # return [opt1, opt2], [sched1, sched2]
```

□ What it does:

- ✓ Centralizes optimizer configuration
- ✓ Supports single/multiple optimizers
- ✓ Automatic learning rate scheduling
- ✓ Automatic **gradient clipping**, **frequency**, etc.



# Putting it altogether ...

Complete Example

```
class LitMNIST(pl.LightningModule):
    # 1. INITIALIZATION
    def __init__(self, hidden_dim=128):
        super().__init__()
        self.save_hyperparameters()
        self.layers = nn.Sequential(
            nn.Linear(28*28, hidden_dim),
            nn.Linear(hidden_dim, 10)
        )

    # 2. TRAIN LOOP
    def training_step(self, batch, batch_idx):
        x, y = batch
        y_hat = self(x)
        loss = F.cross_entropy(y_hat, y)
        self.log("train_loss", loss)
        return loss

    # 3. VALIDATION LOOP
    def validation_step(self, batch, batch_idx):
        x, y = batch
        y_hat = self(x)
        loss = F.cross_entropy(y_hat, y)
        self.log("val_loss", loss)
        return loss

    # 4. TEST LOOP
    def test_step(self, batch, batch_idx):
        x, y = batch
        y_hat = self(x)
        acc = (y_hat.argmax(dim=1) == y).float().mean()
        self.log("test_acc", acc)

    # 5. PREDICTION LOOP
    def predict_step(self, batch, batch_idx):
        x = batch
        return self(x).argmax(dim=1)

    # 6. OPTIMIZERS
    def configure_optimizers(self):
        return torch.optim.Adam(self.parameters(), lr=1e-3)
```

Final Thought: "This isn't a *new way to write PyTorch* → it's just *organizing* what you already write into *logical, reusable components*."



## 3.3 The Trainer

- Once you've organized your PyTorch code into a `LightningModule`,
  - the Trainer automates everything else.

### YOUR CONTROL

You write PyTorch code in `LightningModule`

### AUTOMATED BEST PRACTICES

Trainer executes it optimally using community-proven patterns

You define **WHAT** to compute (forward pass, loss, metrics)

Trainer handles **HOW** to compute it (distribution, optimization, logging)

You can disable **ANY** automation if you need custom behavior

But defaults come from top AI labs

### The Core Philosophy



### 3.3.1 Basic Use

- This is the basic use of the trainer:

#### Basic Use of Trainer

```
model = MyLightningModule()  
  
trainer = Trainer()  
trainer.fit(model, train_dataloader, val_dataloader)
```

# 3.3.1.1 Under the hood ...

## 1. SETUP PHASE

- Moves model to correct device(s)
- Sets up distributed strategy (DDP/FSDP)
- Initializes precision (fp16/bf16)
- Configures all callbacks
- Sets up logging

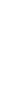


## 2. TRAINING LOOP - AUTOMATED

```
for epoch in range(max_epochs):
    # Training
    for batch in train_dataloader:
        loss = model.training_step(batch)
        loss.backward()          # ← AUTOMATIC
        optimizer.step()         # ← AUTOMATIC
        optimizer.zero_grad()   # ← AUTOMATIC

    # Validation
    if epoch % val_check_interval == 0:
        model.eval()
        with torch.no_grad():
            for batch in val_dataloader:
                model.validation_step(batch)
        model.train()

    # Callbacks execution
    callbacks.on_validation_end()
```



## 4. CLEANUP & RESTORABILITY

- Handles keyboard interrupts (Ctrl+C)
- Can resume from any checkpoint
- Proper distributed teardown



## 3. AUTOMATIC CHECKPOINTING

- Saves best model(s) automatically
- Saves latest model every N epochs
- Manages checkpoint directory

"You write the `training_step()`, Lightning handles **everything else**."



## 3.3.2 Community Wisdom, Built In

Embedded Best Practices from top AI Labs

```
trainer = pl.Trainer()  
# Defaults include:  
# From FAIR's DDP experience:  
# • Proper process group initialization  
# • Correct barrier synchronization  
# • Gradient bucketing optimizations  
  
# From Google's TPU research:  
# • XLA compilation strategies  
# • BF16 mixed precision handling  
# • Distributed checkpointing  
  
# From Stanford's ML research:  
# • Smart validation scheduling  
# • Gradient accumulation patterns  
# • Learning rate warmup strategies  
  
# From MIT's optimization research:  
# • Gradient clipping defaults  
# • Adaptive batch sizing  
# • Numerical stability checks
```

### 3.3.3 Opt – out of Any

#### Opt out of Any

```
trainer = pl.Trainer(  
    # Disable training loop automation  
    automatic_optimization=False, # ← Manual optimization control  
  
    # Manual validation control  
    val_check_interval=None, # ← Run validation manually  
  
    # Custom checkpointing  
    enable_checkpointing=False, # ← Implement your own  
  
    # Disable automatic logging  
    logger=False, # ← Log manually  
)  
# Or override at the LightningModule level:  
class CustomModel(pl.LightningModule):  
    def training_step(self, batch, batch_idx):  
        # Manual optimization example (GANs, Meta-learning)  
        opt1, opt2 = self.optimizers()  
  
        # Compute generator loss  
        self.toggle_optimizer(opt1)  
        loss1 = self.generator_loss(batch)  
        self.manual_backward(loss1)  
        opt1.step()  
        opt1.zero_grad()  
  
        # Compute discriminator loss  
        self.toggle_optimizer(opt2)  
        loss2 = self.discriminator_loss(batch)  
        self.manual_backward(loss2)  
        opt2.step()  
        opt2.zero_grad()
```



### 3.3.4 Real – World Trainer Configurations

One Trainer Infinite Configurations

```
# 1. RESEARCHER ON LAPTOP (Quick Experiments)
researcher = pl.Trainer(
    max_epochs=10,
    fast_dev_run=True,           # Quick sanity check
    overfit_batches=0.01,         # Debug overfitting
    limit_train_batches=0.1,       # Faster iteration
)
```

One Trainer Infinite Configurations

```
# 2. GRAD STUDENT WITH 2 GPUs
grad_student = pl.Trainer(
    accelerator="gpu",
    devices=2,                  # Uses both GPUs automatically
    strategy="ddp",              # Distributed Data Parallel
    precision="16-mixed",        # 2x memory, 1.5x speed
    max_epochs=100,
)
```

"Same code, same **LightningModule**, completely different scale."



### 3.3.4 Real – World Trainer Configurations

One Trainer Infinite Configurations

```
# 3. INDUSTRY TRAINING CLUSTER
production = pl.Trainer(
    accelerator="gpu",
    devices=8,                      # 8 GPUs across nodes
    strategy="fsdp",                 # Fully Sharded Data Parallel
    precision="bf16-mixed",          # Modern precision for large models
    max_epochs=1000,
    gradient_clip_val=1.0,
    accumulate_grad_batches=8,       # Effective batch size = 64
    callbacks=[
        ModelCheckpoint(monitor="val_loss", mode="min"),
        EarlyStopping(monitor="val_loss", patience=20),
        LearningRateFinder(),           # Automatic LR tuning
    ],
    logger=[WandbLogger(), MLFlowLogger()],
    profiler="advanced",             # Performance optimization
)
```

One Trainer Infinite Configurations

```
# 4. GOOGLE TPU POD
tpu_trainer = pl.Trainer(
    accelerator="tpu",
    devices=8,                      # TPU v3-8 pod
    strategy="xla",                 # TPU-specific strategy
    precision="bf16-true",           # Full bfloat16 on TPU
)
```

"Same **code**, same **LightningModule**, completely **different scale**."

# Thank You!!!